

Power BI Processing Patterns

Streamlining Power BI Refreshes Advanced
Methods for Real-Time Refreshes

After this session

Storage modes	Refresh operations	Partitioning	Orchestration
Understand the different storage modes in Power BI and how this relates to refreshes.	Understand which different options there are to refresh your dataset manually as well as in an automated fashion.	Understand how different partitions in the data model are generated and can be triggered to refresh.	Understand how you can orchestrate your Power BI refresh and bind to other processes.



Marc Lelijveld

Solution Architect Data & Analytics
Macaw Netherlands



 @MarcLelijveld

 [linkedin.com/in/MarcLelijveld](https://www.linkedin.com/in/MarcLelijveld)

 Data-Marc.com



What we cover today

- Refreshes?
- Incremental refresh
- High frequency refreshes & hybrid tables
- Refresh individual objects
- Orchestration
- Wrap-up



REFRESH

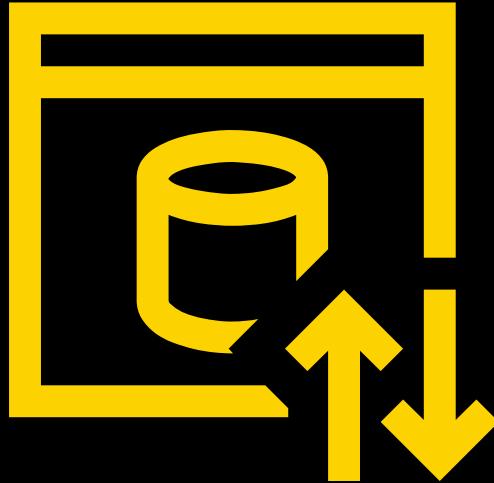


Refreshes

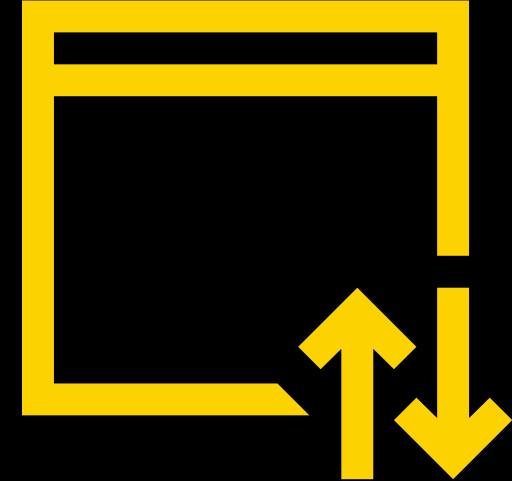
Understanding storage modes



Import



Dual



DirectQuery

Oh wait... we now also have DirectLake!

Storage Modes - Power BI Desktop

File Home Help External Tools

Cut Copy Get Excel Power BI SQL Enter Dataverse Recent sources Data Transform Refresh data Manage relationships Security Q&A setup Language schema Sensitivity Publish Clipboard

Product Subcategory

Product

Internet Sales - Agg

Properties

Name: Internet Sales

Description: Enter a description

Synonyms: Enter a comma-separated list of synonyms for Q&A

Row label: Select a row label

Key column: Select a column with unique values

Is hidden: No

Is featured table: No

Storage mode: DirectQuery

All tables Layout 1 +

Understanding Storage Modes

Three storage modes

- Import – data cached in the model
- DirectQuery – queries are submitted to the back-end data source
- Dual – can act in both above storage modes, depending on query context

Configuring storage modes

- Storage modes are set on table level
- Setting storage mode to Import is an irreversible operation
- Data in DirectQuery mode cannot be displayed in the data tab

Caches and DirectQuery

Risks of mixing storage modes

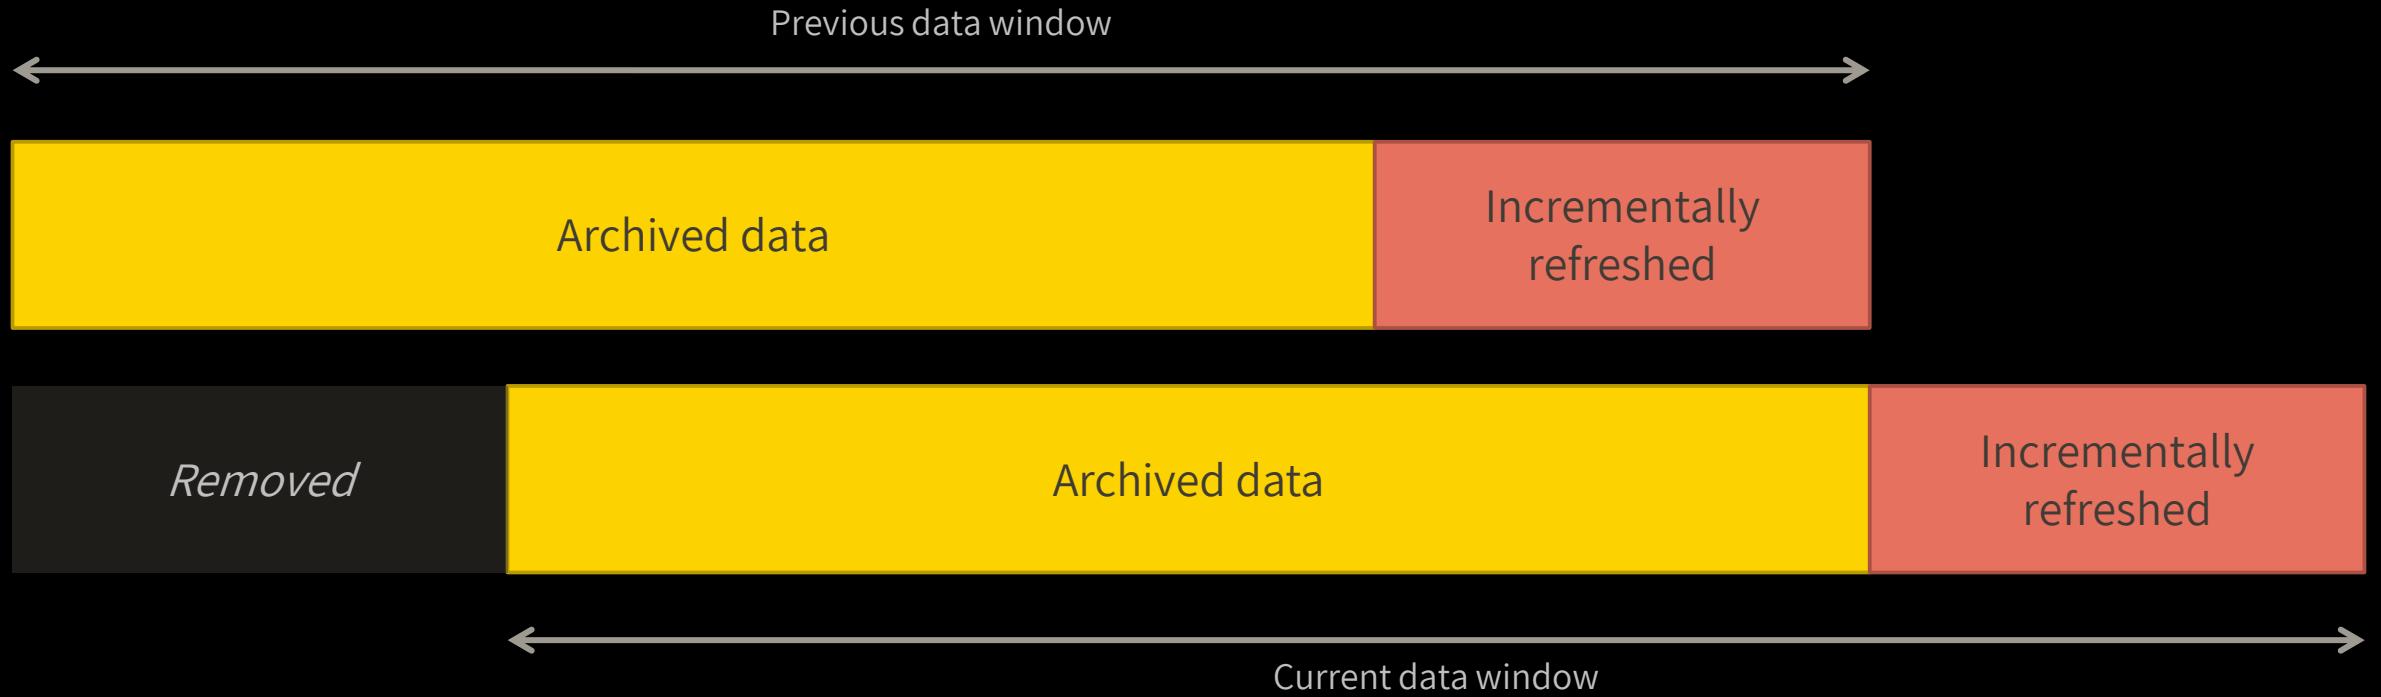
- Avoid mismatch in results when query bits DirectQuery compared to Import
- Data cached (import) could be behind compared to DirectQuery data
- Make sure cached data is kept in sync – regularly refresh!

Benefits to choose one or the other

- Improve query performance – cache data for faster end-user performance
- Data refresh optimization – no need to refresh for non-cached data
- Near-real time requirements – reduce query latency when in DirectQuery mode
- Large datasets – choose to not import certain data

Different refresh operation types

Type	What	Power BI Desktop	Power BI Service
Manual	Individual Table	✓	✗
	Full model	✓	✓
Scheduled	Individual Table	✗	✗
	Full model	✗	✓

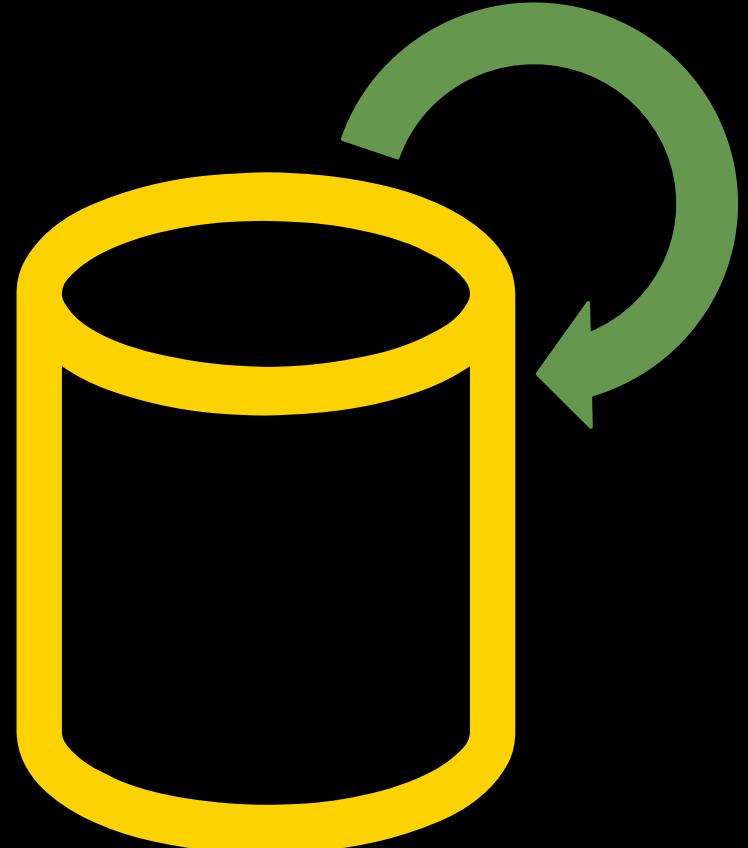


Incremental refresh

Considering incremental refresh

Why should you use it?

- Large data volumes
- Faster refreshes
- Resource consumption reduced
(Both source & Power BI side)



Incremental refresh

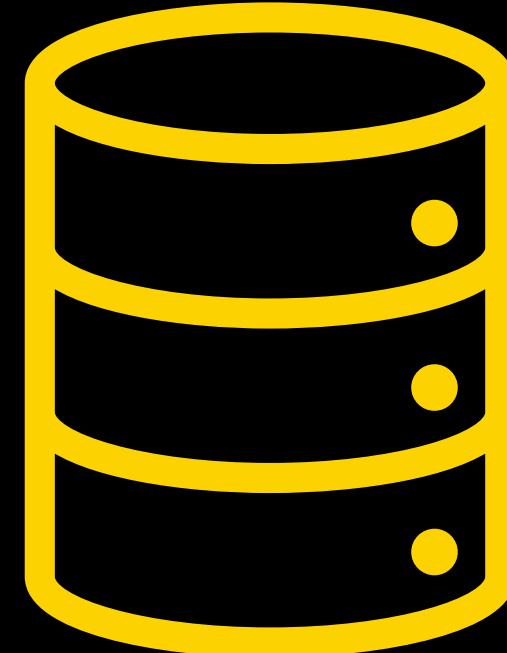
Incremental refresh is supported for Power BI Premium, Premium per user, **Power BI Pro**, and Power BI Embedded datasets.

Getting the latest data in **real time** with DirectQuery is **only supported for Power BI Premium**, Premium per user, and Power BI Embedded datasets.

Supported data sources

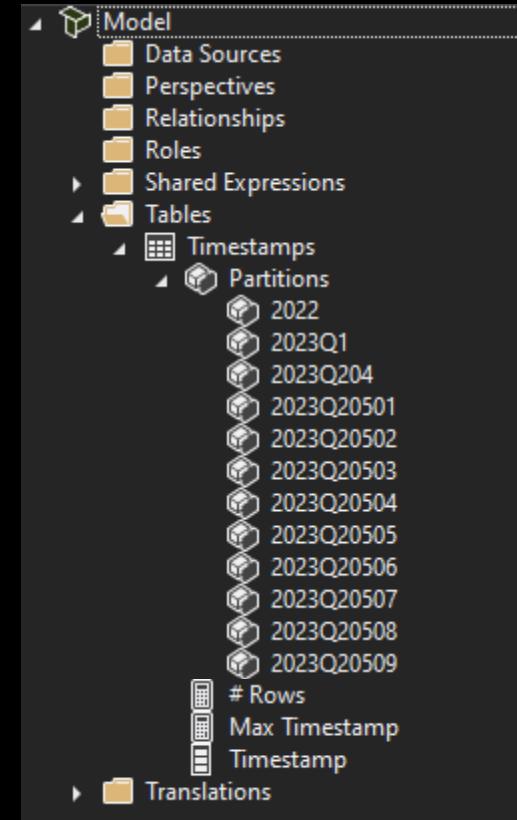
Data sources must be

- Structured
- Relational
- Allow query folding
- Typically – SQL like sources



Breaking up your data in partitions

- Each partition represents a time frame
- New partitions get added automatically
- Old partitions (in archive) get merged



Configure Incremental Refresh

Manage Parameters

New

RangeStart	Name: RangeStart
RangeEnd	Description:
<input checked="" type="checkbox"/> Required	
Type:	Date/Time
Suggested Values:	Any value
Current Value:	1/1/2022 12:00:00 AM

OK Cancel

Incremental refresh and real-time data

Refresh large tables faster with incremental refresh. Plus, get the latest data in real time with DirectQuery (Premium only). [Learn more](#)

① These settings will apply when you publish the dataset to the Power BI service. Once you do that, you won't be able to download it back to Power BI Desktop. [Learn more](#)

1. Select table
Orders

2. Set import and refresh ranges

Incrementally refresh this table

Archive data starting Years before refresh date
Data imported from 1/1/2017 to 11/17/2022 (inclusive)

Incrementally refresh data starting Days before refresh date
Data will be incrementally refreshed from 11/18/2022 to 11/20/2022 (inclusive)

3. Choose optional settings

Get the latest data in real time with DirectQuery (Premium only) [Learn more](#)
Selected table cannot be folded for DirectQuery.

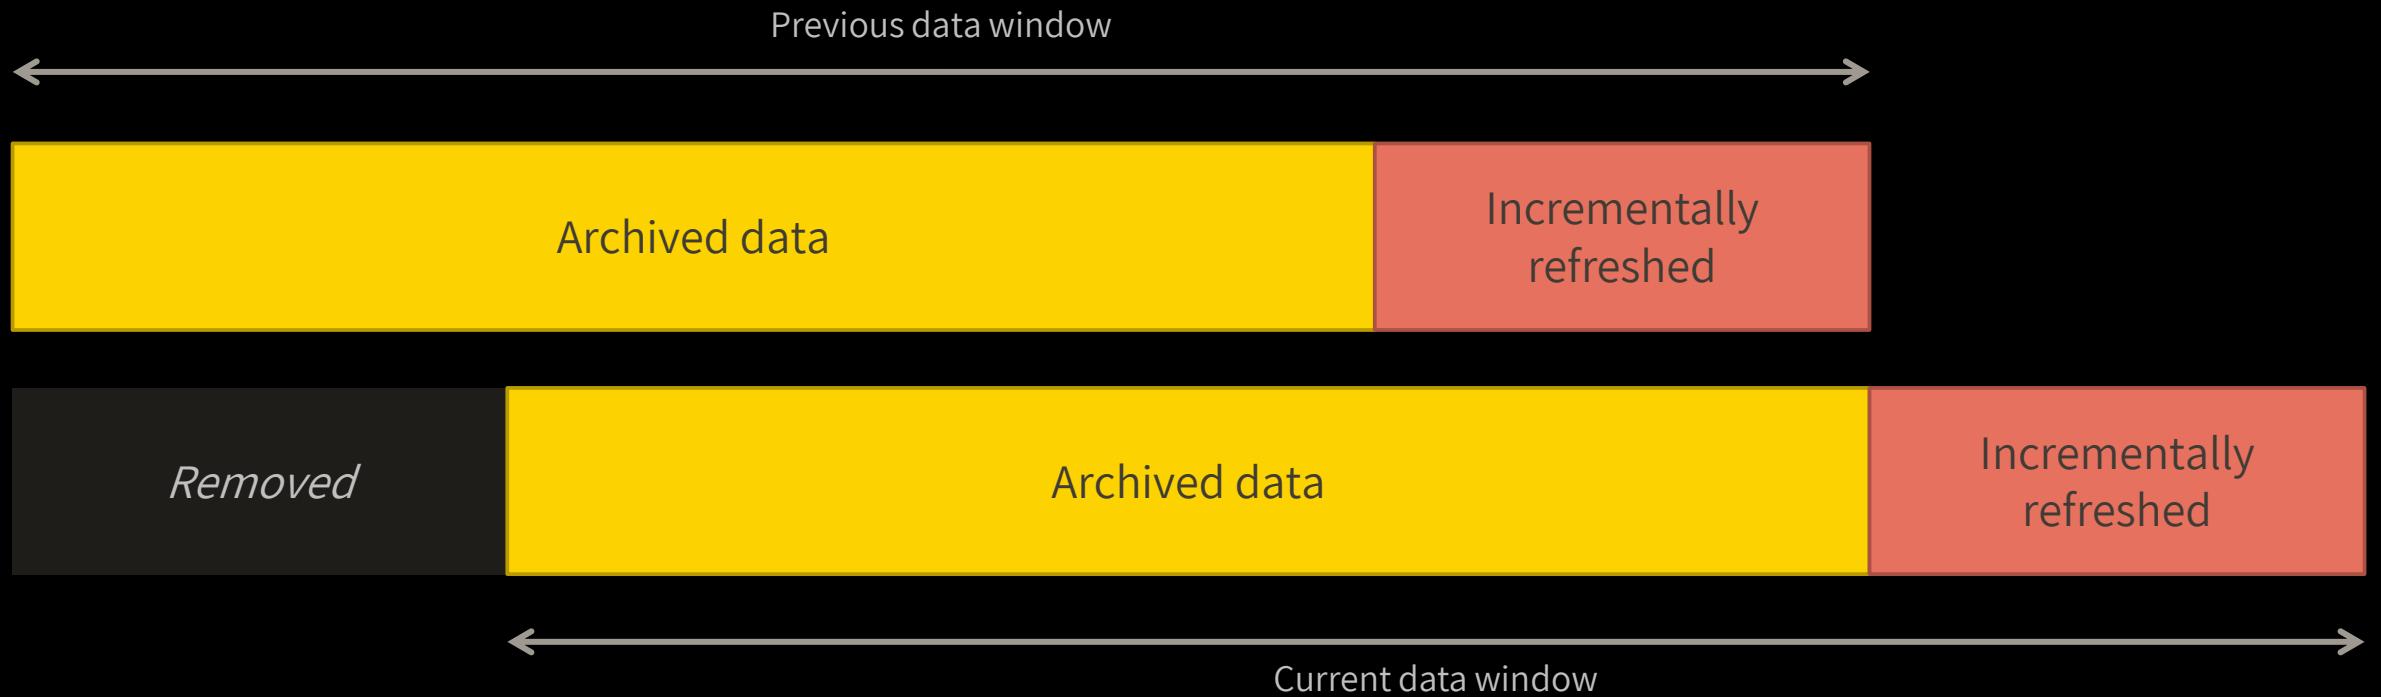
Only refresh complete days [Learn more](#)

Detect data changes [Learn more](#)

4. Review and apply

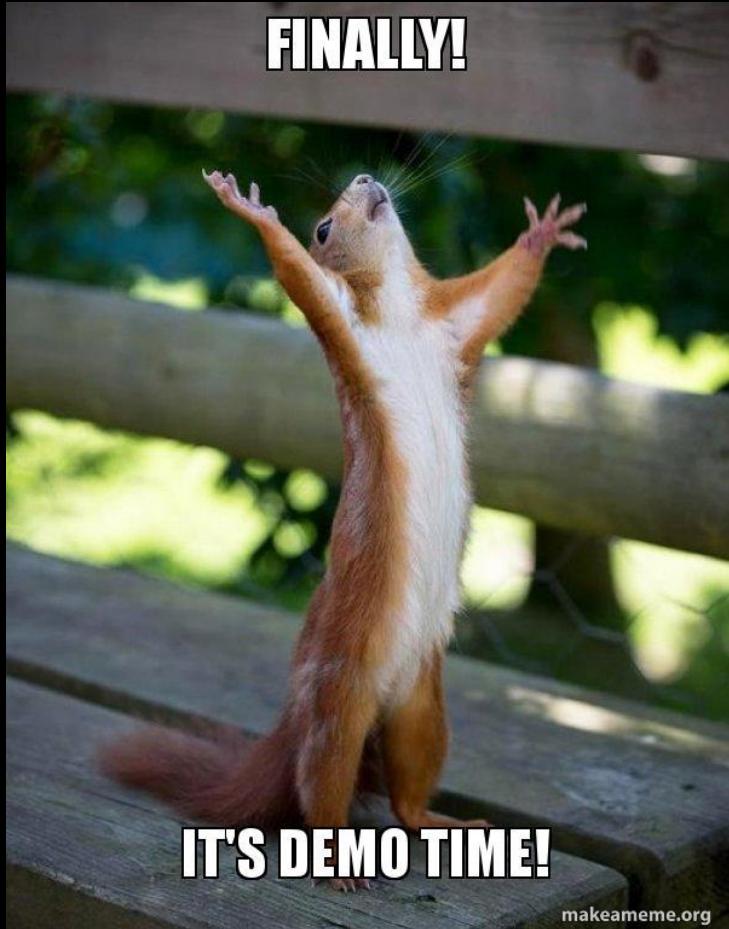
Archived 5 years before Incremental Refresh 3 days before Refresh date

Moving time frame

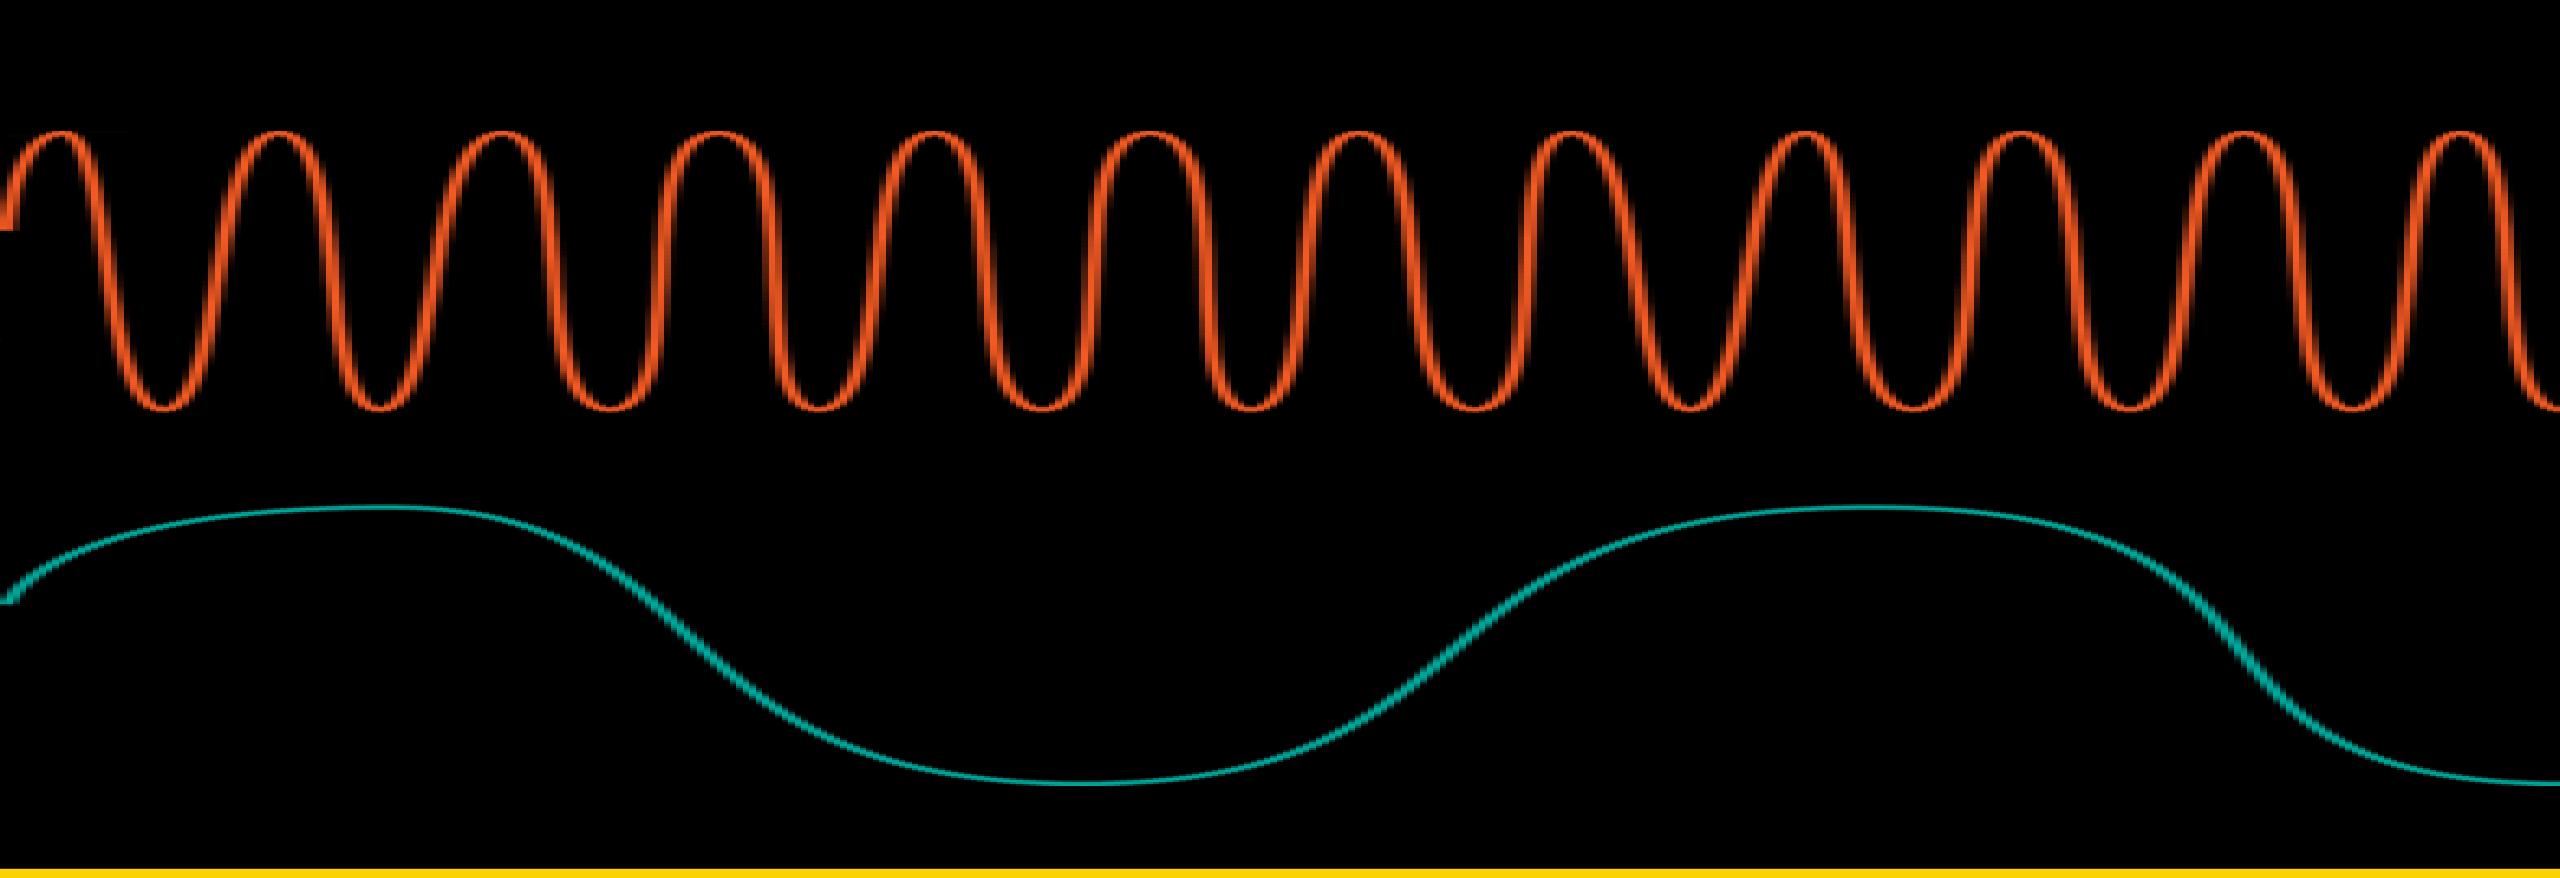


Detect data changes

- Even more selective refresh – only that what has been changed
- Based date/time column used to identify and refresh only those days where the data has changed
- Typically, system generated dates like Modified Date Time
- Should never be the same column as your incremental periods
- The maximum value is evaluated for each partition **in the incremental range** to detect whether the partition should be refreshed or not



Demo - configure
incremental refresh



High frequency refreshes & hybrid tables

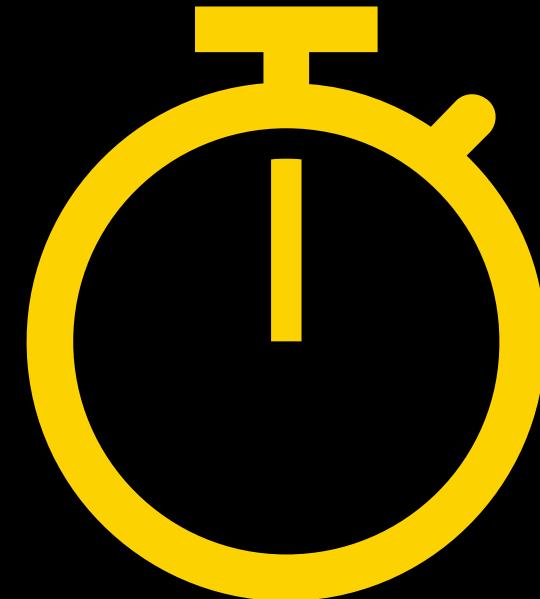
What is high frequency?

Let's say real time?

- Hourly?
- Minutes?
- Seconds?

Or...

- Just daily batches



Refresh limitations

Frequency	Manual	Scheduled (via the service)
Free		Up to 8 times per day
Pro		Up to 8 times per day
Premium per User	Unlimited	48 times per day
Premium Capacity	Unlimited	48 times per day

Hybrid tables

- Live / Realtime data in Power BI
- Combines different storage modes on partition level in a single table
- Goes hand-in-hand with Incremental Refresh

Granularity	Name	Row Count
Year	2011	295,489,717
Year	2012	297,678,498
Year	2013	295,575,442
Year	2014	292,477,875
Year	2015	297,780,469
Year	2016	294,060,081
Year	2017	300,419,682
Year	2018	296,541,108
Year	2019	292,787,420
Year	2020	299,273,979
Quarter	2021Q1	74,135,277
Month	2021Q104	24,939,498
Day	2021Q10501	820,805
Day	2021Q10502	826,885
Day	2021Q10503	821,043
Day-DirectQuery	2021Q10504-DQ	271,110
Total		3,063,898,887

Archived: **Import**

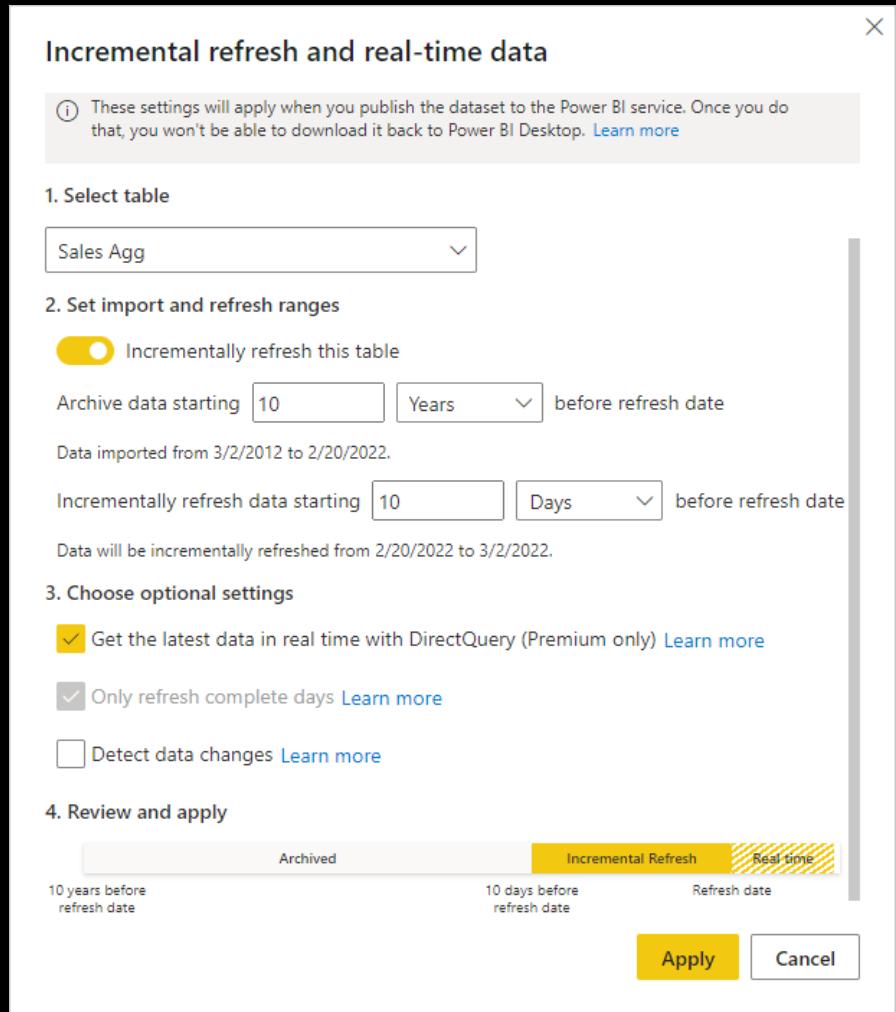
Incremental refresh: **Import**

Real time: **DirectQuery**

Hybrid tables

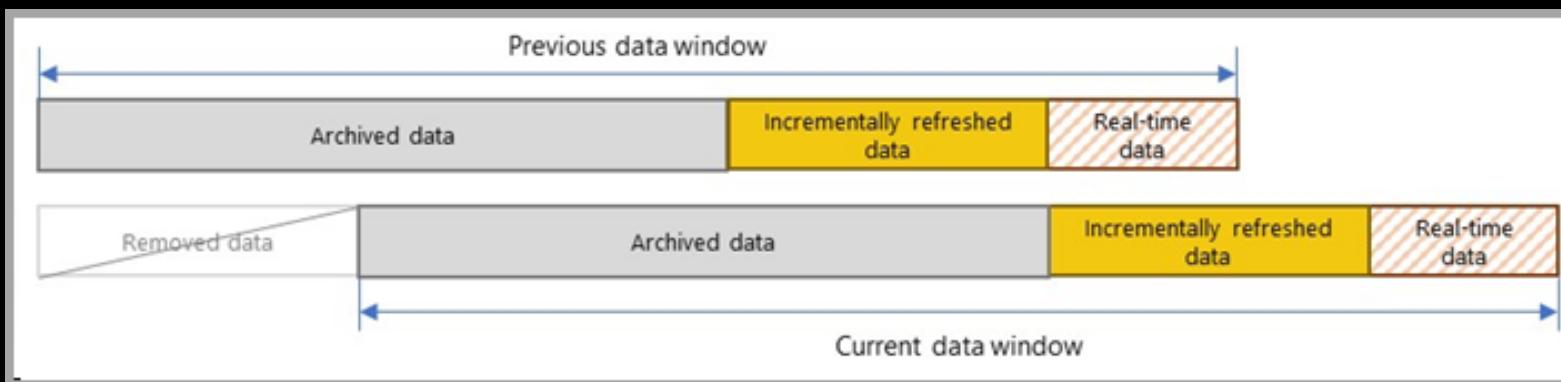
- Implementation with Incremental Refresh
- Customizable via 3rd party tooling like Tabular Editor

Limitation: Only 1 DQ partition per table allowed at the moment.



Hybrid tables – what challenge does it solve?

- Realtime scenarios without full tables on DQ mode
- No complex refresh mechanisms needed with partition refresh and queries over XMLA
- No more multiple tables and complex DAX to combine to achieve the same goal

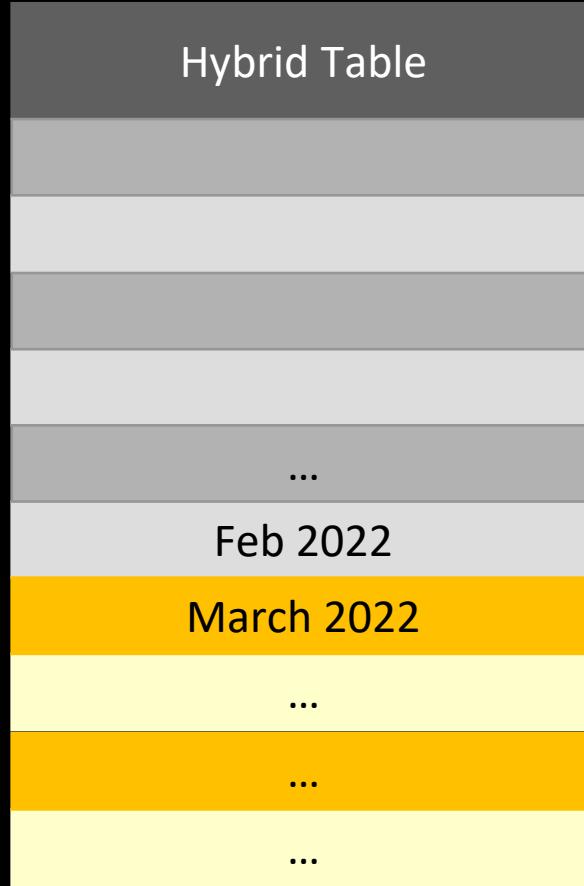


Hybrid tables – Keep in mind that...

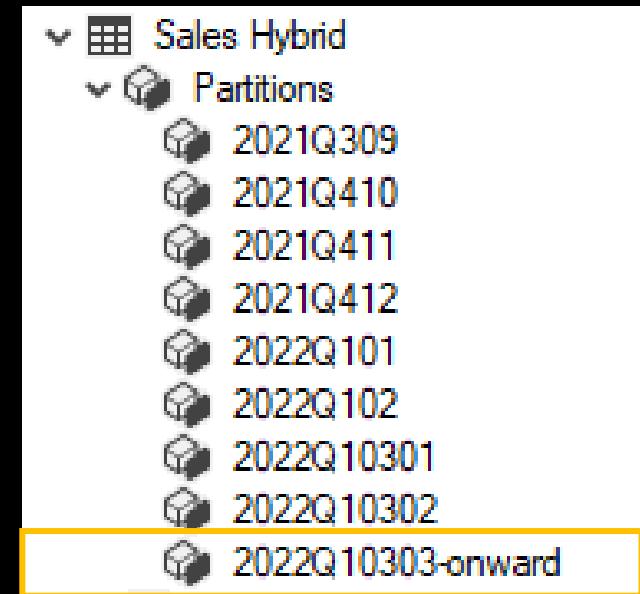
- Premium feature
- DAX restrictions for DirectQuery apply
- Limited Power Query capabilities (due to DQ)
- Requires Large Dataset Format (storage) in workspace
- Performance hit on upstream data sources



Demo - configure
hybrid tables and
streaming data



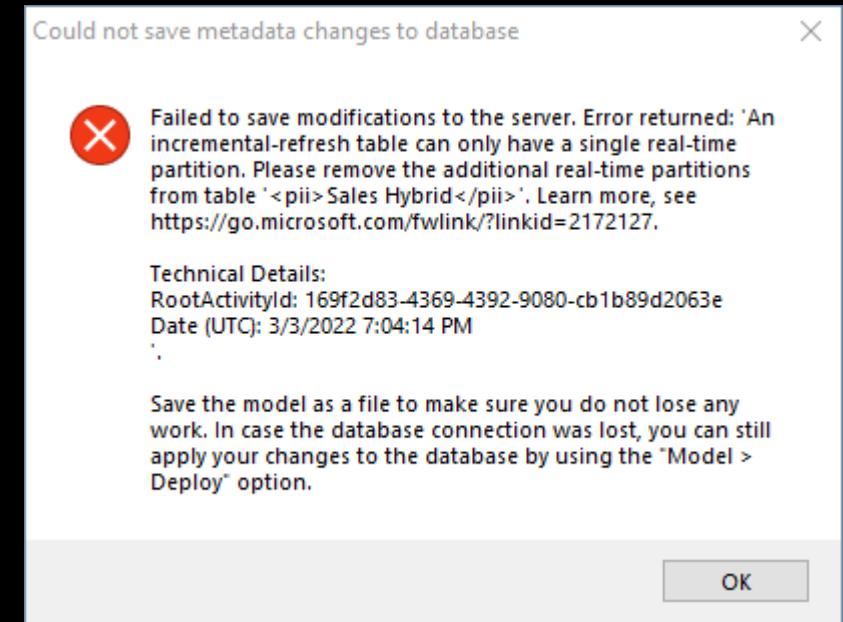
} Import
} DirectQuery →



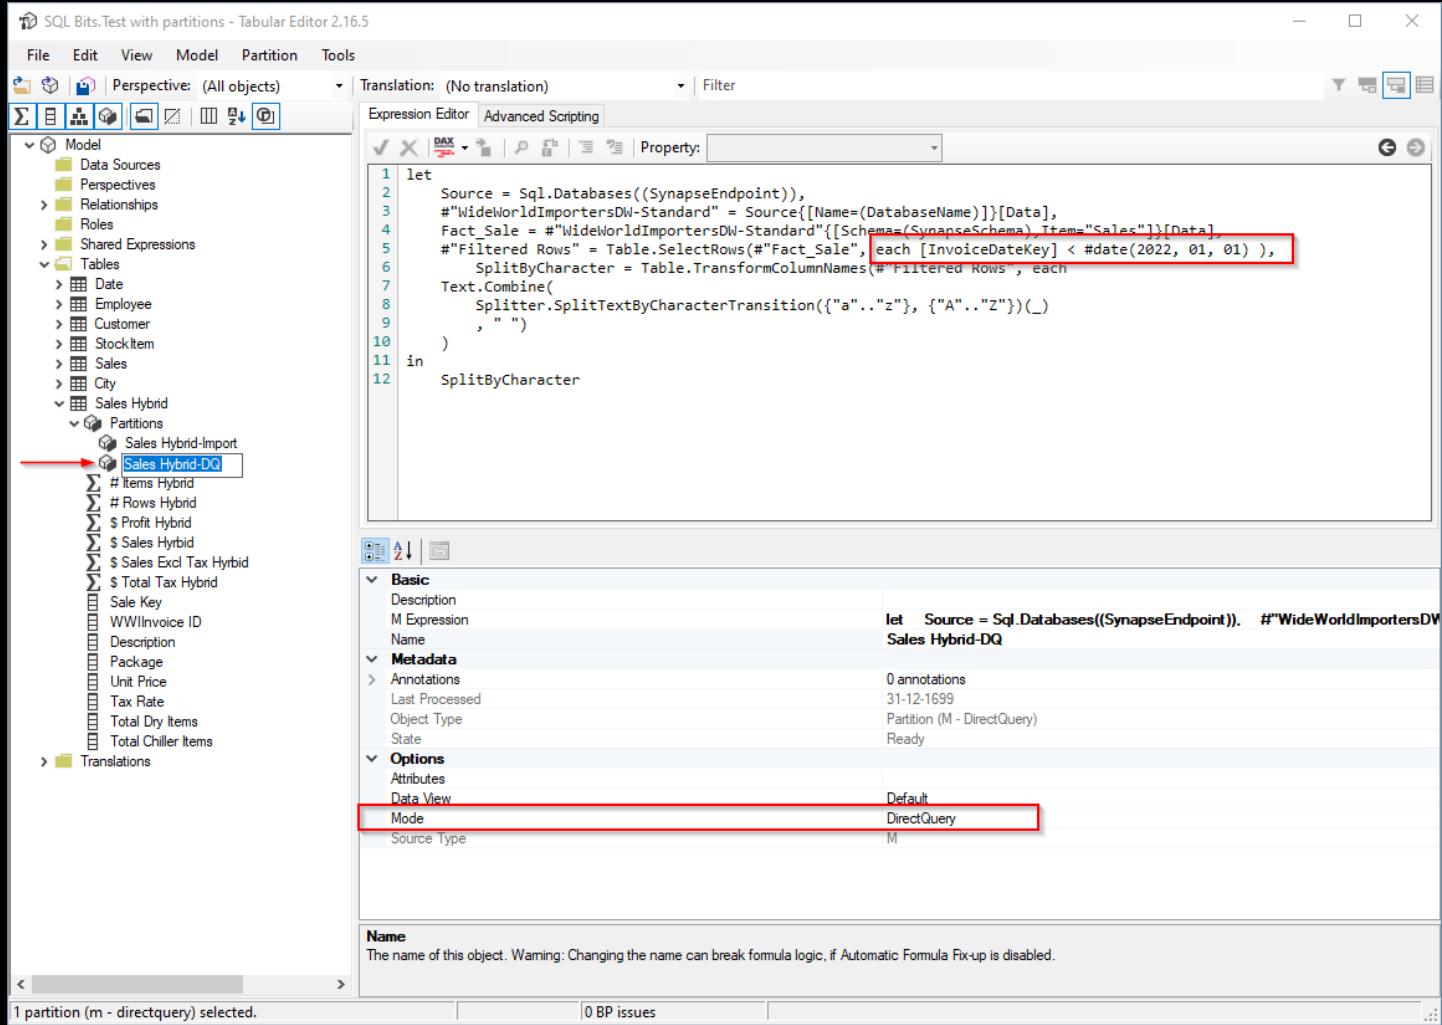
Can I change partition storage modes?

You cannot adjust tables with incremental refresh policies applied. However, there are other options to consider.

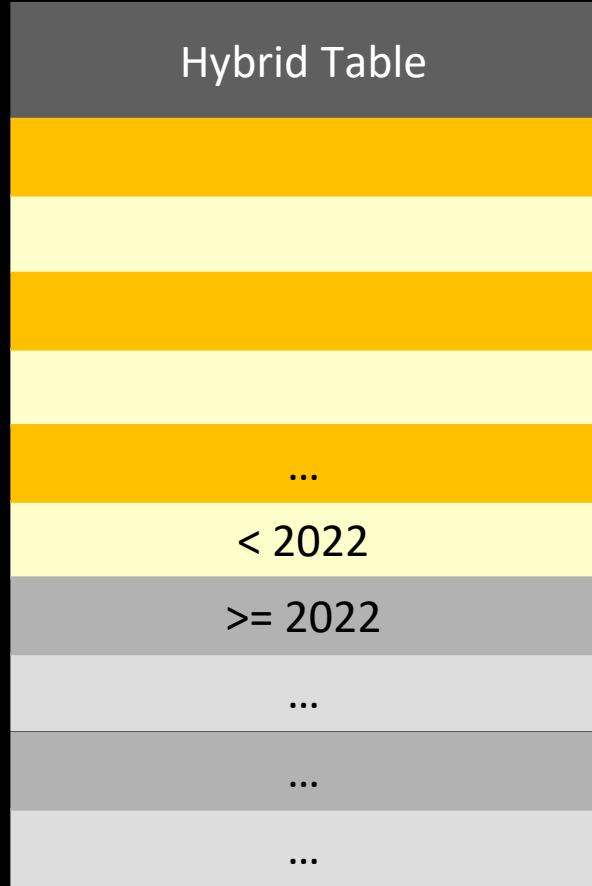
Options	Default
Data View	DirectQuery
Mode	PolicyRange
Source Type	



Manually setup partitioning



Create your own hybrid table setup, using 3rd party tooling like Tabular Editor



} DirectQuery

} Import

DirectQuery Partition definition

```
let
    Source = Sql.Databases((SynapseEndpoint)),
    #"WideWorldImportersDW-Standard" = Source{[Name=(DatabaseName)]}[Data],
    Fact_Sale = #"WideWorldImportersDW-
Standard"[{[Schema=(SynapseSchema),Item="Sales"]}] [Data],
    #"Filtered Rows" = Table.SelectRows(#"Fact_Sale", each
        [InvoiceDateKey] < #date(2022, 01, 01) and
        [InvoiceDateKey] >= #date(2013, 01, 01)),
        SplitByCharacter = Table.TransformColumnNames(#"Filtered Rows", each
            Text.Combine(
                Splitter.SplitTextByCharacterTransition({"a".."z"}, {"A".."Z"})(_)
                , " ")
            )
        in
            SplitByCharacter
```

Import Partition definition

```
let
    Source = Sql.Databases((SynapseEndpoint)),
    #"WideWorldImportersDW-Standard" = Source{[Name=(DatabaseName)]}[Data],
    Fact_Sale = #"WideWorldImportersDW-
Standard"[{[Schema=(SynapseSchema),Item="Sales"]}] [Data],
    #"Filtered Rows" = Table.SelectRows(#"Fact_Sale", each
        [InvoiceDateKey] >= #date(2022, 01, 01) and
        [InvoiceDateKey] < #date(2022, 03, 31)),
        SplitByCharacter = Table.TransformColumnNames(#"Filtered Rows", each
            Text.Combine(
                Splitter.SplitTextByCharacterTransition({"a".."z"}, {"A".."Z"})(_)
                , " ")
            )
        in
            SplitByCharacter
```

Downsides

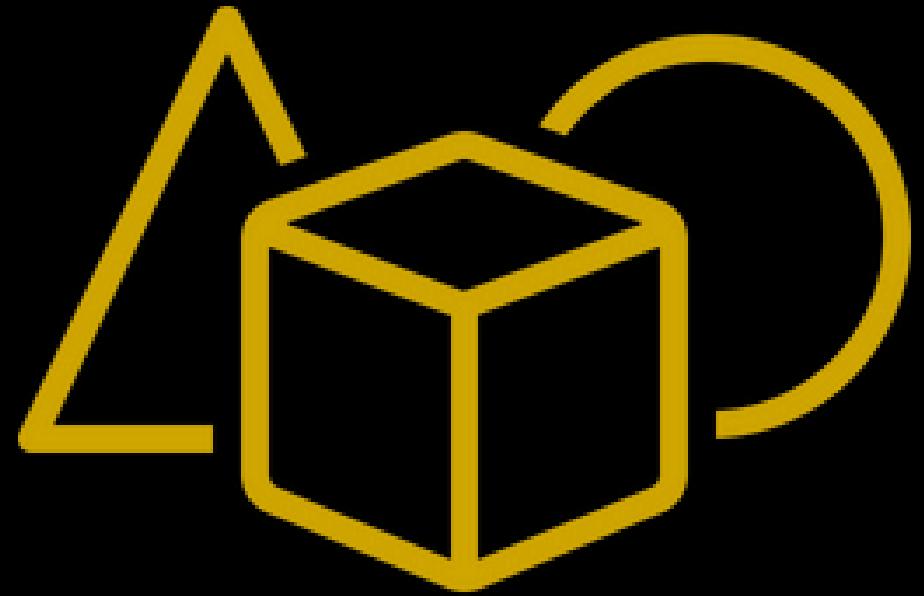
Potential performance impact

- Policy definitions and partitions defined by Power BI will be scanned more optimally when a DAX query is executed
- With manual setup partitions, your DAX expression might be broken up in multiple sub-queries which hit all partitions



Caution

At this time, a write operation on a dataset authored in Power BI Desktop will prevent it from being downloaded back as a PBIX file. Be sure to retain your original PBIX file.



Refresh individual objects

Effective refreshing

Considerations

- Refreshing the entire model takes too long with high load on sources
- Can we only refresh certain tables?
- Can we only refresh certain partitions?
- Can we use DQ tables/partitions (Hybrid Tables)

What do we need?

- Enhanced refresh API
- XMLA Endpoints

Smaller increments



Bind your refresh operations in Power BI to refresh operations in for example your data platform



Run smaller increments to speed up the end-to-end process



Trigger one (or multiple) selective table(s) or partition(s) at a time.

Smaller increments

Smaller increments can only be triggered to refresh via:

- Manual processing via Management Studio (or other tools)
- Power BI Enhanced Refresh API - Previously known as Async refresh API
- XMLA endpoints

Enhanced refresh API

More granular controls

- type - full / clearValues / calculate / dataonly / automatic / defragment
- Commitmode – default is transactional
- maxParallism – 10 max
- retryCount – number of retries before failing
- objects – which objects to refresh
- applyRefreshPolicy – true/false whether you want to apply potential incremental refresh if configured
- effectiveDate – in case of incremental refresh, this parameter overwrites the current date

Enhanced refresh API

Specify the objects to refresh

POST

https://api.powerbi.com/v1.0/myorg/groups/f089354e-8366-4e18-aea3-4cb4a3a50b48/datasets/cfafbeb1-8037-4d0c-896e-a46fb27ff229/refreshes

```
{  
  "type": "Full",  
  "commitMode": "transactional",  
  "maxParallelism": 2,  
  "retryCount": 2,  
  "objects": [  
    {  
      "table": "DimCustomer",  
      "partition": "DimCustomer"  
    },  
    {  
      "table": "DimDate"  
    }  
  ]  
}
```



Demo – Refresh one table or partition

More than refreshes

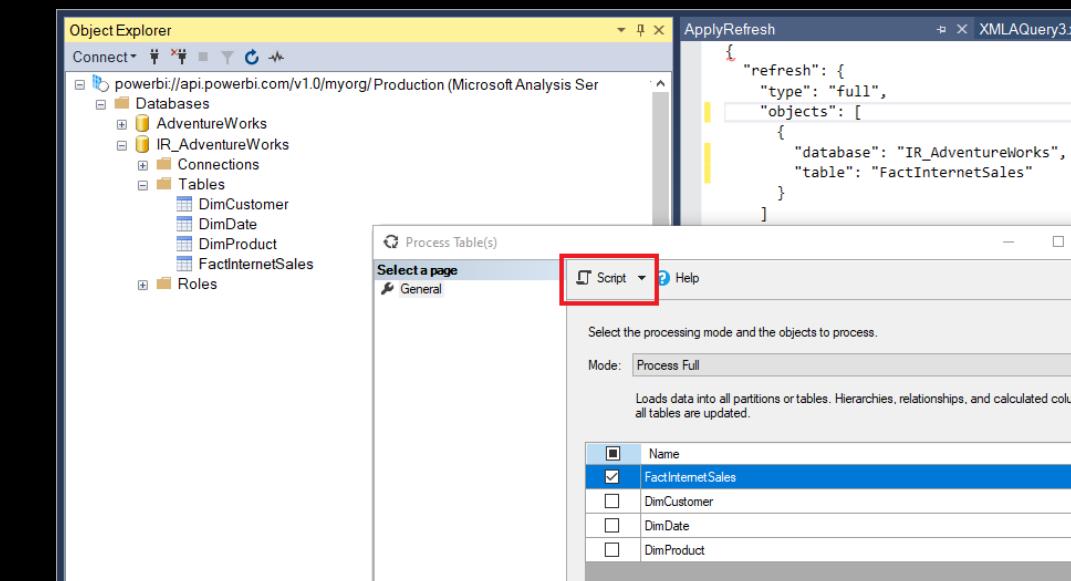
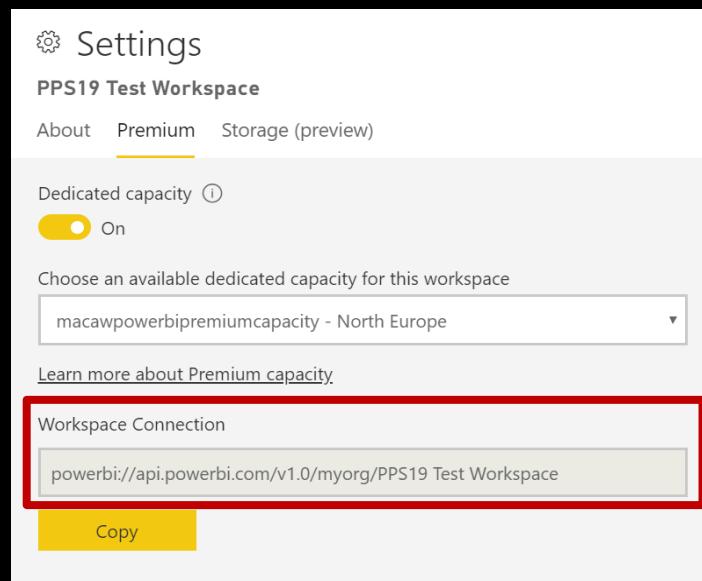
Enhanced refresh API also allows you to

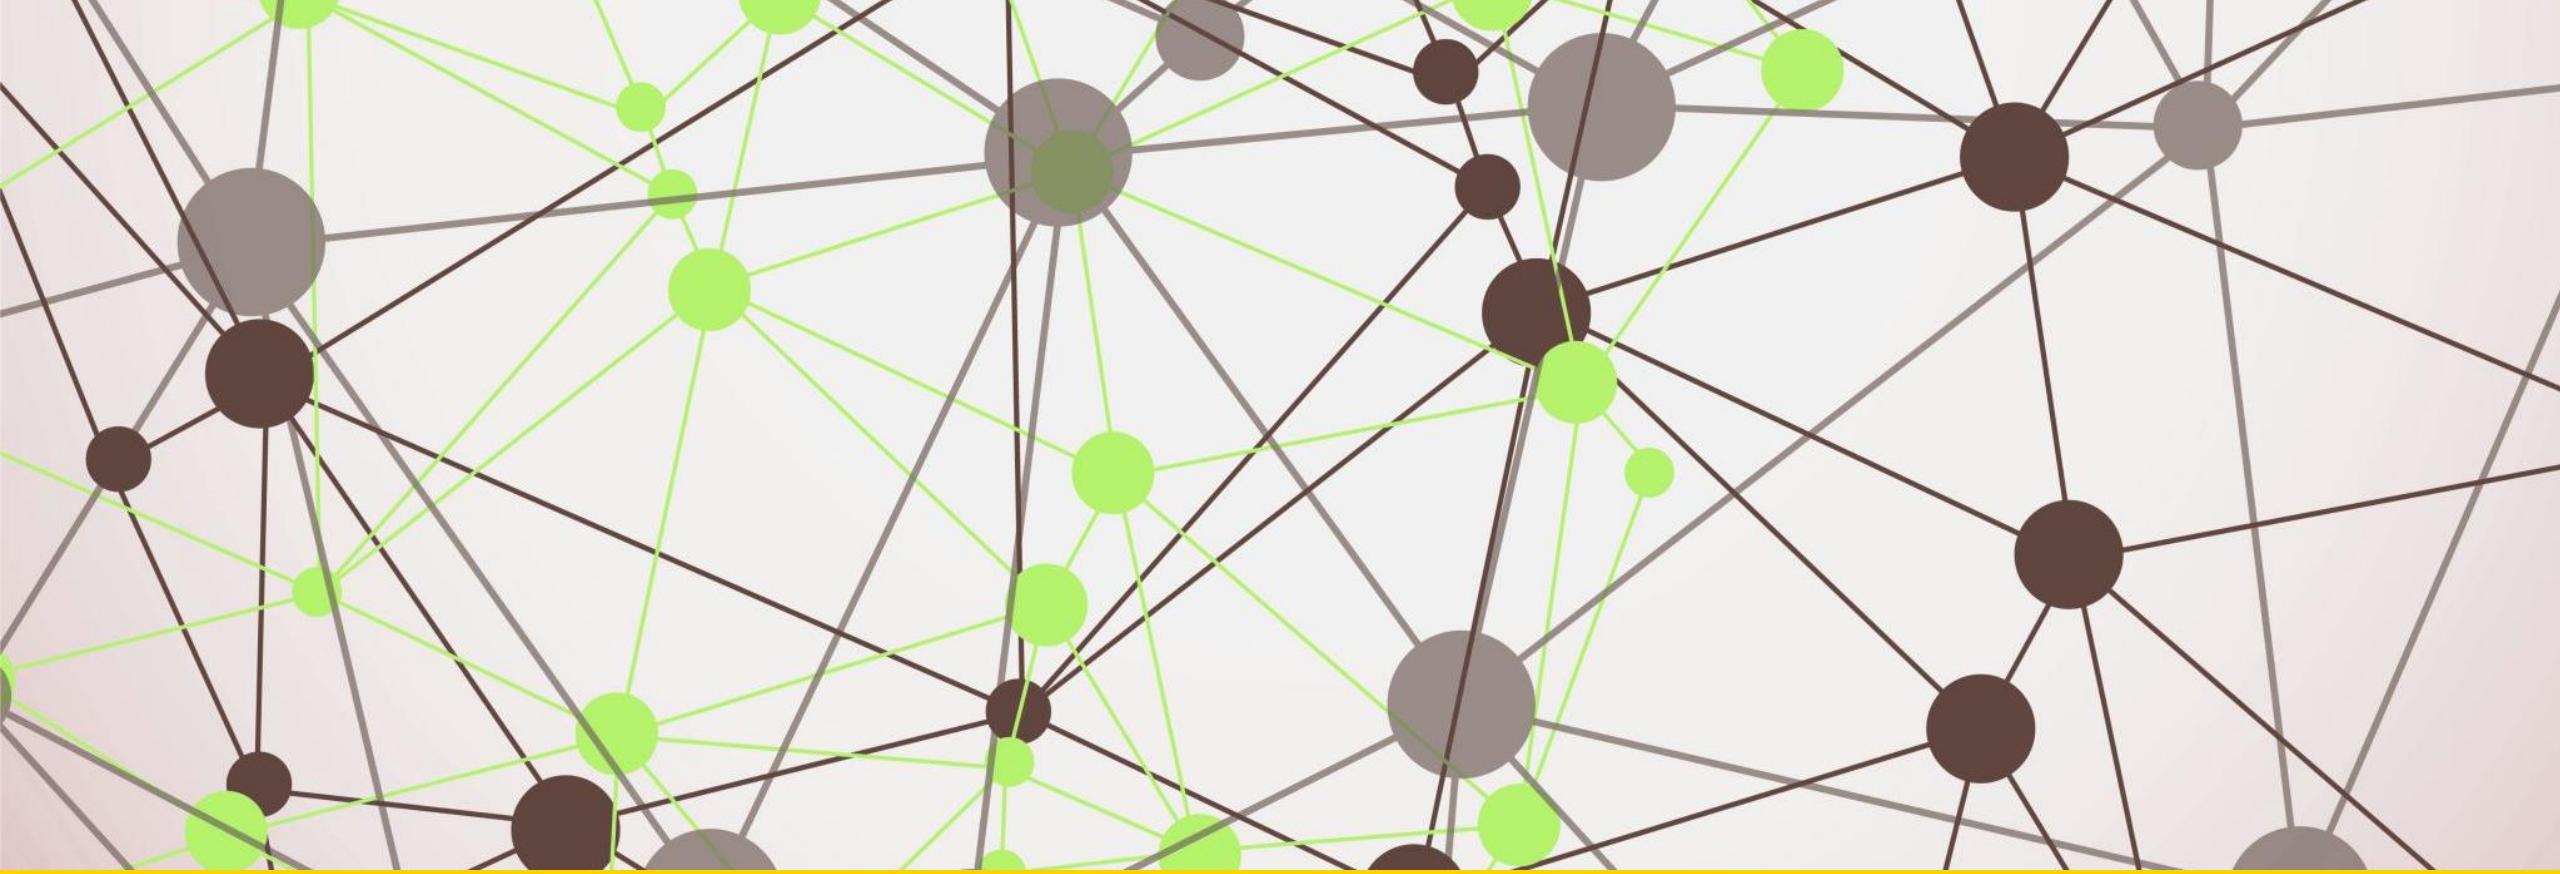
- **GET** refresh history
- **GET** individual refresh details like progress or failures
- **DELETE** a running refresh

```
JSON
{
  "startTime": "2020-12-07T02:06:57.1838734Z",
  "endTime": "2020-12-07T02:07:00.4929675Z",
  "type": "Full",
  "status": "InProgress",
  "currentRefreshType": "Full",
  "objects": [
    {
      "table": "DimCustomer",
      "partition": "DimCustomer",
      "status": "InProgress"
    },
    {
      "table": "DimDate",
      "partition": "DimDate",
      "status": "InProgress"
    }
  ]
}
```

Is the API the only option?

- Execute operations over XMLA endpoints
- Using Tabular Model Scripting Language
- More granular control using SQL Service Management Studio



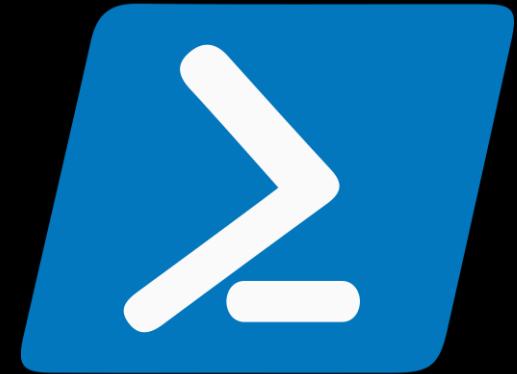
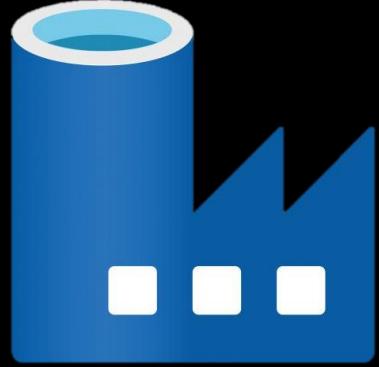


Binding refreshes to processes

Central E2E orchestration

- Combine pipelines from Data Platform with Power BI
- Lowest latency between source and report
- Incremental loading where possible
- Consider including backup operations for Power BI

What to use to trigger and bind together?





Demo – Refresh API from Synapse Pipeline

Wrap up

LET'S
RECAP...

Refreshes can be manual or **automated**

Incremental refresh makes your refresh **more efficient**

With **detect data changes**, further optimization can be done

Partitioning in incremental refresh should be **based on functional dates**

Lower latency by linking Power BI refreshes to other processes

Use the enhanced refresh API to **trigger individual objects**

Use the toolkit of your preference for **end-to-end orchestration**

Resources

Setting up scheduled refreshes

<https://learn.microsoft.com/en-us/power-bi/connect-data/refresh-scheduled-refresh>

Configure incremental refresh

<https://learn.microsoft.com/en-us/power-bi/connect-data/incremental-refresh-configure>

Enhanced refresh API

<https://learn.microsoft.com/en-us/power-bi/connect-data/asynchronous-refresh>

Dataset connectivity with XMLA endpoints

<https://learn.microsoft.com/en-us/power-bi/enterprise/service-premium-connect-tools>

Tabular Model Scripting Language (TMSL)

<https://learn.microsoft.com/en-us/analysis-services/tmsl/tabular-model-scripting-language-tmsl-reference?view=asallproducts-allversions>

These slides

<https://github.com/marcelijveld/Slide-decks>

Thanks for attending!



Marc Lelijveld

Solution Architect Data & Analytics
Macaw Netherlands



@MarcLelijveld

[linkedin.com/in/MarcLelijveld](https://www.linkedin.com/in/MarcLelijveld)

Data-Marc.com

